

Support Vector Machines

In this exercise sheet, you will experiment with training various support vector machines on a subset of the MNIST dataset composed of digits 5 and 6. First, download the MNIST dataset from <http://yann.lecun.com/exdb/mnist/>, uncompress the downloaded files, and place them in a `data/` subfolder. Install the optimization library CVXOPT (python-cvxopt package, or directly from the website www.cvxopt.org). This library will be used to optimize the dual SVM in part A.

Part A: Kernel SVM and Optimization in the Dual

We would like to learn a nonlinear SVM by optimizing its dual. An advantage of the dual SVM compared to the primal SVM is that it allows to use nonlinear kernels such as the Gaussian kernel, that we define as:

$$k(x, x') = \exp\left(-\frac{\|x - x'\|^2}{\sigma^2}\right)$$

The dual SVM consists of solving the following quadratic program:

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{ij} \alpha_i \alpha_j y_i y_j k(x_i, x_j)$$

subject to:

$$0 \leq \alpha_i \leq C \quad \text{and} \quad \sum_{i=1}^N \alpha_i y_i = 0.$$

Then, given the alphas, the prediction of the SVM can be obtained as:

$$f(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^N \alpha_i y_i k(x, x_i) + \theta > 0 \\ -1 & \text{if } \sum_{i=1}^N \alpha_i y_i k(x, x_i) + \theta < 0 \end{cases}$$

where

$$\theta = \frac{1}{\#SV} \sum_{i \in SV} \left(y_i - \sum_{j=1}^N \alpha_j y_j k(x_i, x_j) \right)$$

and SV is the set of indices corresponding to the unbound support vectors.

Implementation (25 P)

We will solve the dual SVM applied to the MNIST dataset using the CVXOPT quadratic optimizer. For this, we have to build the data structures (vectors and matrices) to must be passed to the optimizer.

- *Implement* a function `gaussianKernel` that returns for a Gaussian kernel of scale σ , the Gram matrix of the two data sets given as argument.
- *Implement* a function `getQPMatrices` that builds the matrices P, q, G, h, A, b (of type `cvxopt.matrix`) that need to be passed as argument to the optimizer `cvxopt.solvers.qp`.
- *Run* the code below using the functions that you just implemented. (It should take less than 3 minutes.)

```
In [1]: import utils, numpy, cvxopt, cvxopt.solvers
import solutions

Xtrain, Ttrain, Xtest, Ttest = utils.getMnist56()
```

```

cvxopt.solvers.options['show_progress'] = False

for scale in [10,30,100]:
    for C in [1,10,100]:

        # Prepare kernel matrices

        ### TODO: REPLACE BY YOUR OWN CODE
        Ktrain = solutions.gaussianKernel(Xtrain,Xtrain,scale)
        Ktest  = solutions.gaussianKernel(Xtest,Xtrain,scale)
        ###

        # Prepare the matrices for the quadratic program

        ### TODO: REPLACE BY YOUR OWN CODE
        P,q,G,h,A,b = solutions.getQPMatrices(Ktrain,Ttrain,C)
        ###

        # Train the model (i.e. compute the alphas)
        alpha = numpy.array(cvxopt.solvers.qp(P,q,G,h,A,b)['x']).flatten()

        # Get predictions for the training and test set
        SV = (alpha>1e-6)
        uSV = SV*(alpha<C-1e-6)
        theta = 1.0/sum(uSV)*(Ttrain[uSV]-numpy.dot(Ktrain[uSV,:],alpha*Ttrain)).sum()
        Ytrain = numpy.sign(numpy.dot(Ktrain[:,SV],alpha[SV]*Ttrain[SV])+theta)
        Ytest  = numpy.sign(numpy.dot(Ktest[:,SV],alpha[SV]*Ttrain[SV])+theta)

        # Print accuracy and number of support vectors
        Atrain = (Ytrain==Ttrain).mean()
        Atest  = (Ytest ==Ttest ).mean()
        print('Scale=%3d C=%3d SV: %4d Train: %.3f Test: %.3f'%(scale,C,sum(SV),Atrain,Atest))
    print('')

Scale= 10 C= 1 SV: 1000 Train: 1.000 Test: 0.937
Scale= 10 C= 10 SV: 1000 Train: 1.000 Test: 0.937
Scale= 10 C=100 SV: 1000 Train: 1.000 Test: 0.937

Scale= 30 C= 1 SV: 254 Train: 1.000 Test: 0.985
Scale= 30 C= 10 SV: 274 Train: 1.000 Test: 0.986
Scale= 30 C=100 SV: 256 Train: 1.000 Test: 0.986

Scale=100 C= 1 SV: 317 Train: 0.973 Test: 0.971
Scale=100 C= 10 SV: 159 Train: 0.990 Test: 0.975
Scale=100 C=100 SV: 136 Train: 1.000 Test: 0.975

```

Analysis (10 P)

- Explain which combinations of parameters σ and C lead to good generalization, underfitting or overfitting?
- Explain which combinations of parameters σ and C produce the fastest classifiers (in terms of amount of computation needed at prediction time)?

Part B: Linear SVMs and Gradient Descent in the Primal

The quadratic problem of the dual SVM does not scale well with the number of data points. For large number of data points, it is generally more appropriate to optimize the SVM in the primal. The primal optimization problem for linear SVMs can be written as

$$\min_{w, \theta} ||w||^2 + C \sum_{i=1}^N \xi_i \quad \text{where} \quad \forall_{i=1}^N : y_i(w \cdot x_i + \theta) \geq 1 - \xi_i \quad \text{and} \quad \xi_i \geq 0.$$

It is common to incorporate the constraints directly into the objective and then minimizing the unconstrained objective

$$J(w, \theta) = ||w||^2 + C \sum_{i=1}^N \max(0, 1 - y_i(w \cdot x_i + \theta))$$

using simple gradient descent.

Implementation (15 P)

- *Implement* the function J computing the objective $J(w, \theta)$
- *Implement* the function DJ computing the gradient of the objective $J(w, \theta)$ with respect to the parameters w and θ .
- *Run* the code below using the functions that you just implemented. (It should take less than 1 minute.)

```
In [2]: import utils, numpy
import solutions

C = 10.0
lr = 0.001

Xtrain, Ttrain, Xtest, Ttest = utils.getMnist56()

n, d = Xtrain.shape

w = numpy.zeros([d])
theta = 1e-9

for it in range(0, 101):

    # Monitor the training and test error every 5 iterations
    if it%5==0:
        Ytrain = numpy.sign(numpy.dot(Xtrain, w) + theta)
        Ytest = numpy.sign(numpy.dot(Xtest, w) + theta)

        ### TODO: REPLACE BY YOUR OWN CODE
        Obj = solutions.J(w, theta, C, Xtrain, Ttrain)
        ###

        Etrain = (Ytrain == Ttrain).mean()
        Etest = (Ytest == Ttest).mean()
        print('It=%3d J: %9.3f Train: %.3f Test: %.3f' % (it, Obj, Etrain, Etest))

    ### TODO: REPLACE BY YOUR OWN CODE
    dw, dtheta = solutions.DJ(w, theta, C, Xtrain, Ttrain)
    ###
```

```
w = w - lr*dw
theta = theta - lr*dtheta
```

It= 0	J: 10000.000	Train: 0.471	Test: 0.482
It= 5	J: 68520.417	Train: 0.961	Test: 0.958
It= 10	J: 49918.674	Train: 0.973	Test: 0.961
It= 15	J: 37473.229	Train: 0.973	Test: 0.963
It= 20	J: 28590.129	Train: 0.974	Test: 0.965
It= 25	J: 21746.877	Train: 0.977	Test: 0.967
It= 30	J: 16987.200	Train: 0.980	Test: 0.968
It= 35	J: 13646.095	Train: 0.986	Test: 0.967
It= 40	J: 11187.127	Train: 0.986	Test: 0.967
It= 45	J: 9182.940	Train: 0.991	Test: 0.967
It= 50	J: 7692.273	Train: 0.990	Test: 0.968
It= 55	J: 6437.609	Train: 0.988	Test: 0.966
It= 60	J: 5253.071	Train: 0.995	Test: 0.966
It= 65	J: 4515.520	Train: 0.992	Test: 0.967
It= 70	J: 4016.851	Train: 0.996	Test: 0.966
It= 75	J: 3647.983	Train: 0.997	Test: 0.965
It= 80	J: 3497.204	Train: 0.998	Test: 0.966
It= 85	J: 3404.280	Train: 1.000	Test: 0.966
It= 90	J: 3336.804	Train: 1.000	Test: 0.966
It= 95	J: 3270.665	Train: 1.000	Test: 0.966
It=100	J: 3205.837	Train: 1.000	Test: 0.966